

The AI Infrastructure Protocol

July 17, 2026

CONTENTS

name: ai-infrastructure-protocol description: > The AI Infrastructure Protocol — a gated methodology + skill suite that takes a company

The AI Infrastructure Protocol

The five pillars

The journey — five phases, five gates

G0 — SCAN (pillar ② opens here)

G1 — BLUEPRINT (pillar ② closes here)

G2 — BUILD (pillars ①③ execute here)

G3 — SECURE (pillar ⑤ — the hard gate)

G4 — INVESTOR-READY (pillar ④)

Deliverable shape (what the buyer receives)

Pillar → skill map

Operating rules


name: ai-infrastructure-protocol description: > The AI Infrastructure Protocol — a gated methodology + skill suite that takes a company from “we barely use AI” to a fully AI-run, secure, investor-attractive operation. Five pillars (App, Infrastructure, Agents, Security, Investor) driven through one gated pipeline. Use when starting a company AI-infrastructure engagement, running a company scan, building agents/automations for a business, hardening an AI build, or packaging a company for investors.

The AI Infrastructure Protocol

You are the operator of a build engine that installs **maximum AI involvement** into a real company — and proves, at each gate, that what you installed is real, safe, and investor-grade. This document is the spine. Each pillar has its own skill; this file governs how they fit together and in what order they run.


The promise to the buyer: *"You want a top-tier AI-run company but don't know AI. We scan everything you do, then build the agents, automations, apps, knowledge, and systems that run it — secured, and packaged so investors want in."*

The five pillars



Pillar ⑤ (Security) is **not a phase** — it is a gate that every buildable artifact (app, agent, automation, API) must pass before it ships. Pillars ①–④ are the journey; ⑤ is the standard they are held to.

The journey — five phases, five gates



Gate discipline mirrors the app engine: a gate is a check that writes `.protocol/gates/gate-N.json`; the next phase refuses to start until the prior gate is green; a green gate goes stale on drift. Interventions are logged. The `INFRASTRUCTURE-REPORT.md` is the handover artifact — the company's new AI operating system, on paper, with proof.

G0 — SCAN (PILLAR ② OPENS HERE)

AI reads the whole company: every process, tool, data source, workflow, handoff, and pain. Output: `company-scan.json` — the as-is operating map (who does what, with which tools, on what data, where the hours and errors go). This is not a survey; it is a structured audit. The scan skill (`pillars/2-infrastructure`) runs it. **Gate 0 green** = the scan covers every department/process the company named, each with its tools, data, volume, and pain quantified.

G1 — BLUEPRINT (PILLAR ② CLOSES HERE)

From the scan, AI designs the **target AI operating model**: for every process, the decision is *keep human / augment with AI / fully automate*, and the blueprint names the specific agent, automation, app, knowledge base, or workflow that delivers it. Output: `blueprint.json` — the to-be map plus a ranked build backlog (impact × effort). **Gate 1 green** = every process from the scan has a disposition, every automation has a named owner artifact, and the backlog is ranked with an AI-involvement score.

G2 — BUILD (PILLARS ①③ EXECUTE HERE)

Work the backlog. Apps go through the **four-gate app engine** (`pillars/1-app` → `3day-protocol/`). Agents and automations go through the **agent factory** (`pillars/3-agents`). Each built artifact registers in `.protocol/artifacts.json`. **Gate 2 green** = every backlog item marked "this sprint" is built, registered, and individually green on its own sub-gates (an app run's own SHIP-REPORT green, or an agent factory `verify.json` green).

Two gate numberings — do not confuse them. The app engine (`3day-protocol/`) runs its OWN P0–P5 / gate-0..3 machine inside each app's `~/builds/<artifactId>/.protocol/`, where *its* gate-2 is security and *its* gate-3 is deploy-verify. The protocol (this file) runs G0–G4 in the **engagement-root** `.protocol/`, where **G3 = SECURE (pillar ⑤)**. They live in separate directories and never overwrite each other. When this document says “G3” it always means the protocol’s SECURE gate, never the app engine’s deploy-verify gate.

G3 — SECURE (PILLAR ⑤ — THE HARD GATE)

Every artifact from G2 is hardened and proven: secrets management, auth + least- privilege, RLS/data-access, API rate-limiting + abuse protection, input validation, dependency/secret scanning. **Nothing ships that Security hasn’t seen.** This reuses the app engine’s security gate for apps and adds agent/API-specific checks. **Gate 3 green** = every registered artifact passes the security skill’s checklist; a missing tool or an unproven check is RED, never a skip.

G4 — INVESTOR-READY (PILLAR ④)

Package the now-AI-run company so an investor says “*I need a stake.*” AI-involvement becomes the story: the moat (proprietary agents/automations), the unit economics (cost taken out, throughput added), the defensibility, the scalability. Output: `investor-package/` — metrics dashboard, architecture narrative, unit-economics model, and the one-line thesis. **Awaits Ian’s investor documentation + online research to finalize the scoring rubric.** **Gate 4 green** = the package answers every criterion in the investor rubric with a real number or artifact from this build, not a claim.

Deliverable shape (what the buyer receives)

The protocol is delivered as **knowledge (PDFs) + skills**:

- `knowledge/` — PDF instruction + knowledge documents (generated via the `make-pdf` skill): the operating manual, the pillar playbooks, the security standard, the investor rubric. These are what a human reads.
- `skills/` — the executable Claude Code skills (this build copied 36 of Ian’s existing skills, mapped to pillars, plus the new pillar skills). These are what the operator runs.
- `pillars/` — each pillar’s own SKILL.md governing its phase.

Pillar → skill map

Pillar	New skill	Draws on (copied skills)
① App	(uses <code>3day-protocol/</code>)	build/: supabase, deploy-to-vercel, interface-design, ian-design-standards, web-design-guidelines, shadcn
② Infrastructure	<code>pillars/2-infrastructure/SKILL.md</code> (scan→blueprint)	agents/: agent-browser, scrape, composio-cli, dispatching-parallel-agents
③ Agents	<code>pillars/3-agents/SKILL.md</code> (agent factory)	agents/: composio-cli, whatsapp-ai-agent, hook-generator, voice-builder, pair-agent
④ Investor	<code>pillars/4-investor/SKILL.md</code> (investability)	business/: enotna-ekonomika, vrednostno-cenovanje, distribucijski-kanali, pozicioniranje, pricing, marketing-plan, revops, competitor-profiling; advisors/: hormozi, christensen, dunford, skok, campbell, kim-mauborgne, marks
⑤ Security	<code>pillars/5-security/SKILL.md</code> (hardening gate)	security/: cso, guard, careful

Operating rules

- Maximize AI involvement is the objective function.** At every blueprint decision, the default is “can AI do this?” Human-only is the exception that must be justified, not the baseline.
- Never claim, always prove.** Every gate produces evidence. The INFRASTRUCTURE-REPORT cites artifacts, not adjectives.
- Security is a gate, not a feature.** No artifact reaches the investor package without passing G3.

4. **Reuse before build.** 36 skills are already bundled; the operator composes them before writing anything new.
5. **The company's words beat the operator's assumptions.** The scan records what the company actually does, in their language, before any AI redesign.