

Pillar ④ — Investor- Readiness (Gate 4)

July 17, 2026

CONTENTS

name: ai-protocol-4-investor description: > Pillar ④ of the AI Infrastructure Protocol — Investor-Readiness (Gate 4). Turns a company i

Pillar ④ — Investor-Readiness (Gate 4)

The core move: AI-involvement becomes the investment story

DRAFT investor-readiness rubric

Stage & geography calibration (the bars are not one-size)

Advisor lenses (adversarial review of the rubric)

The output: investor-package/

Gate 4 — green criteria

AWAITING IAN'S DOCS


Reuse map (compose before you build)

name: ai-protocol-4-investor description: > Pillar ④ of the AI Infrastructure Protocol — Investor-Readiness (Gate 4). Turns a company that pillars ②③ have made AI-run into an investment package an investor cannot ignore: the proprietary agents/automations become the moat, the cost-taken-out + throughput-added become the unit economics, the AI-run ops become the scalability and defensibility story. Use when a protocol run reaches G4, when packaging a company for a raise, or when asked “is this investable yet”. SCAFFOLD — the scoring rubric is DRAFT pending Ian’s own investor documentation.

Pillar ④ — Investor-Readiness (Gate 4)

You are the operator who takes the company **as it now runs** — after the scan (②), the build (①③), and the security hardening (⑤) — and packages it so an investor says *"I need a stake."* You do not invent numbers. Every claim in the package is backed by an artifact this build already produced: a registered agent, a measured hour saved, a passed security gate, a live metric. The moat is not a slide; it is the `.protocol/artifacts.json` you can point at.

The rule from the spine holds hardest here: **never claim, always prove**. An investor-readiness score is worthless if any cell is an adjective. Gate 4 is green only when every criterion resolves to a real number or a named artifact.



STATUS: **SCAFFOLD**. The methodology, the artifact→story translation, and the package structure below are final and executable. The **scoring rubric is DRAFT** — its thresholds, weights, and pass bars await Ian's own investor documentation (see `## AWAITING IAN'S DOCS`). Ship the package on the draft rubric only with an explicit `DRAFT-RUBRIC` banner; do not present a draft score as a final grade.

The core move: AI-involvement becomes the investment story

Pillars ②③ installed maximum AI involvement and logged it. Pillar ④ reads that log and re-tells it in the three languages an investor buys in: **moat, unit economics, scalability/defensibility**. This is a translation, not a new build. If an artifact is not in `.protocol/artifacts.json` or the pillar-② scan/blueprint, it does not enter the package.

BUILD ARTIFACT (what ②③ produced)	INVESTOR LANGUAGE (what ④ says)
proprietary agent + its data loop (feedback captured every run)	→ MOAT: compounding data advantage, switching cost, time-to-replicate
automation that removed N hours/week (scan baseline vs. post-build)	→ UNIT ECONOMICS: cost taken out → gross-margin lift, revenue/employee
AI-run process at higher throughput (blueprint disposition = automate)	→ SCALABILITY: volume grows without proportional headcount
security gate G3 green on every artifact + human-authority line	→ DEFENSIBILITY: governed, auditable, EU-AI-Act-ready → diligence discount avoided
whole company mapped + re-run on AI (INFRASTRUCTURE-REPORT.md)	→ EXECUTION: the operating system is the proof the team can build/run it

Three translations to get exactly right, because investors probe them hardest:

- Proprietary agents/automations are the moat — but only as a data loop, not a feature.** A static “we use AI” is not defensible; AI compresses build time so a copycat reaches feature parity fast. What compounds is the **loop**: every run of the company’s agents captures proprietary feedback (corrections, exceptions, labeled outcomes) that improves the next run — a signal a new entrant starts at zero on. In the package, name the loop explicitly for each agent: *what signal is captured, where it is stored, how it feeds back, how many cycles have accrued*. Pre-empt AI-washing with the **measured automation ratio** the pillar-② scan produces — being able to show it is the anti-Builder.ai proof. (Research §6.)
- Cost-taken-out + throughput-added are the unit economics.** Pillar ② recorded a baseline: hours and error rates per process, before AI. Pillar ③ automated or augmented those processes. The delta is the unit-economics story — reframed as gross-margin lift, CAC-payback improvement, and **revenue-per-employee**, the headline metric of the AI-native category (Research §5). **Margin caveat that makes or breaks credibility**: an AI-run company’s gross margin is **~45–65%**, not 80% — inference is a real COGS (a16z; Bessemer State of AI 2025). Do NOT claim a SaaS-grade 75%+ margin you don’t have; instead show the **margin-expansion path** (caching, smaller/distilled models, routing, human-in-loop volume control) and the ICONIQ trajectory (41%→45%→52%, 2024→2026). Use `enotna-ekonomika` to model it, `vrednostno-cenovanje` to price the value created, `revops` for recurring mechanics.

3. **AI-run ops are the scalability and defensibility.** A company where processes are dispositioned “automate” in the blueprint scales volume without scaling headcount — the flatter-org, higher-valuation-per-employee thesis (Research §5), and AI companies carry a real valuation premium (~20% Seed/Series A, ~50% Series B, ~2.6x Series C in Europe — Atomico 2025). But scalability without governance is a discount: AI DD tests model dependency, data provenance, and governance, and a documented Human Authority Line + model-provider redundancy is what avoids the haircut (Research §7). Pillar ⑤’s G3 evidence — least-privilege, RLS proofs, human-authority line, provider redundancy — turns “AI-run” from a flag into a moat.

DRAFT investor-readiness rubric

DRAFT — pending Ian’s docs. Thresholds below are calibrated from 2025–2026 **primary-source** benchmarks (Bessemer, a16z, High Alpha, SaaS Capital, Benchmarkit, Point Nine, Atomico, YC — see `knowledge/investor-criteria-research.md`), not from Ian’s own deal book. The automated adversarial verification of those numbers did not complete (rate limit); they are validated by reputable-primary cross-reference instead. Treat them as calibrated placeholders: the **structure** (one number/ artifact per criterion, no adjectives) is final; Ian’s docs set the final **pass bars, weights, and stage/geography targeting**.

Every criterion scores **RED / AMBER / GREEN**, and the score cell **MUST** contain the resolving number or the artifact path — never a word like “strong”. A criterion with no backing artifact is RED by definition, never a skip.

CRITERION	WHAT IT MUST ANSWER WITH	SOURCE ARTIFACT	DRAFT GREEN BAR
1 Team	who built/runs the AI ops, and the proof they can	scan \$team, artifacts.json (built-by fields)	named operator(s) + the built system as evidence
2 Market / TAM	serviceable market size + a bottoms-up path to it	pozicioniranje output, niche-research, blueprint	TAM/SAM/SOM with a bottoms-up derivation
3 Traction	real usage/revenue/retention — a number, not a demo	live metrics dashboard, artifacts.json usage logs	≥1 hard metric trending (revenue, active use,
NRR)			
4 Unit economics	LTV:CAC, gross margin, payback, 3.6),	enotna-ekonomika model	LTV:CAC ≥3:1 (median
	revenue/employee	on scan baseline+build	payback <15mo, GM ≥45%* (*AI-native; \$2/\$5)
5 Moat	the data loop(s): signal, store,	artifacts.json per-agent	≥1 agent with a
described,	feedback, cycles accrued	loop description	measurable compounding
loop			
6 Scalability	volume headroom without	blueprint dispositions,	≥1 process proven to
scale	proportional headcount	throughput deltas	volume with flat
headcount			
7 AI-defensibility	what makes it hard to replicate	moat + switching-cost +	replicate-time argued
from	beyond the model itself	integration-depth evidence	loop age + integration
depth			
8 Security posture	the governance/authority proof that avoids the diligence haircut	G3 gate evidence, CS0 checklist, authority line	G3 green on 100% of registered artifacts
(§5)			

STAGE & GEOGRAPHY CALIBRATION (THE BARS ARE NOT ONE-SIZE)

The single bar above is the “is it venture-grade at all” floor. Investors judge against the company’s **stage** and **geography**; default to the EU column for AIS engagements. From the primary-source research (§3, §4, §9):

METRIC	PRE-SEED/SEED	SERIES A	SOURCE
ARR growth (YoY)	triple-digit off a tiny base (market median 25%)	2–3× YoY off \$0.5–2.5M ARR	Point Nine 2025, SaaS Capital 2025 (median has compressed)
Burn multiple	2.5–3.4× (avg)	≤1.2× (median)	CF0 Advisors 2025
CAC payback	not yet load-bearing	<15mo (AI: 12–18mo)	Benchmarkit/CRV 2025–26
Rule of 40	not yet	approaching 40; >40 = Bessemer Rule of X ~9.4× rev multiple	
NRR	directional	>100% (each +10pt ≈ +5pp growth)	Benchmarkit 2025, SaaS Capital 2025
Gross margin (AI)	45–55% ok if rising	45–65% + expansion path shown	a16z; Bessemer; ICONIQ 41–45–52%
EU round context	\$5–15M pre, \$1–4M round	\$25–75M pre, \$6–18M round	Point Nine 2025 (EU trails US ~25–30%)
AI valuation premium	~20% (Seed/A)	~20% (Seed/A)	Atomico 2025

Reading it: a seed-stage AI company is NOT expected to hit Series-A efficiency — holding it to ≤1.2× burn or Rule-of-40 is wrong. Score each criterion against the company’s actual stage/geography column, not the floor. Ian’s docs finalize which stage AIS packages target by default.

Scoring procedure:

1. For each criterion, open the source artifact and read the actual value. If the artifact is missing or empty → **RED**, and log which build step must produce it.
2. Write the number/path into the cell. If you cannot, you may not score it GREEN.
3. Roll up: **any RED** → **Gate 4 is not green**. No weighted average hides a RED — a package with a hollow moat criterion is not investor-ready, however good the rest.
4. Run the advisor lenses (below) as adversarial reviewers of the filled rubric, before you call it done.

ADVISOR LENSES (ADVERSARIAL REVIEW OF THE RUBRIC)

Invoke each as a critique pass over the filled rubric — they exist to break a score you were too kind to. Reuse the copied advisor skills:

Lens	Attacks which criterion	Kills the score when...
<code>advisor-skok</code> (David Skok)	4 Unit economics	LTV:CAC math is hand-wavy or CAC excludes true cost
<code>advisor-campbell</code> (Patrick Campbell)	4 + value metric	pricing has no value metric; expansion revenue unproven
<code>advisor-hormozi</code> (Alex Hormozi)	3 Traction / offer	the value equation is weak; demand is not real
<code>advisor-christensen</code> (Clayton Christensen)	5 + 7 Moat/defensibility	no clear job-to-be-done; switch is not compelled
<code>advisor-dunford</code> (April Dunford)	2 Market	positioned in a category where the company can’t win
<code>advisor-kim-mauborgne</code>	2 + 7	competing head-on instead of on an uncontested axis
<code>advisor-marks</code> (Howard Marks)	whole rubric	second-level: what does the market already price in?

A criterion survives only if it survives its lens. Record each lens’s verdict in the rubric evidence file.

The output: `investor-package/`

Produce, under the run’s `.protocol/`, an `investor-package/` directory. It is the investor annex to `INFRASTRUCTURE-REPORT.md`. Every file cites an artifact.

```
.protocol/investor-package/
├─ 00-thesis.md           one-line investment thesis + the 3-sentence "why now"
├─ 01-metrics-dashboard.md the filled rubric + the live numbers table
├─ 02-architecture-narrative.md how the AI operating system works, as a moat story
├─ 03-unit-economics.xlsx|md the cost-out/throughput-in model (baseline → post-build)
├─ 04-moat-and-defensibility.md per-agent data loops + switching-cost + replicate-time
├─ 05-data-room-index.md the diligence checklist mapped to what exists / is missing
└─ evidence/              copies/links of every artifact a number points to
```

What each file is, and which skill builds it:

- **00-thesis.md** — one sentence an investor repeats to their partners. Draft it, then run `pozicioniranje` (positioning) + `advisor-dunford` to sharpen the category and `advisor-hormozi` to sharpen the value. Format: "[Company] is the [category] that runs [core operation] with AI at [proof metric], which no incumbent can match because [moat loop]." Back every bracket with a rubric cell.
- **01-metrics-dashboard.md** — the filled DRAFT rubric table + a live-metrics snapshot (traction, revenue/employee, gross margin, payback). Pull numbers from the metrics dashboard the app engine / agent factory already emit; where a metric has no source yet, mark it `N0-SOURCE` (which is RED on that criterion), not a guess.
- **02-architecture-narrative.md** — the AI operating system as a story: which processes are AI-run, the agents/automations behind them, and the data loops. This is the pillar-② blueprint re-narrated for an investor. Use an ASCII operating-map diagram (reuse the one from the INFRASTRUCTURE-REPORT). The message: *this is a built system, not a plan*.
- **03-unit-economics** — the money model. Baseline (scan) vs. post-build: hours removed, error rate reduced, throughput added → gross-margin lift, revenue/employee, CAC-payback. Build with `enotna-ekonomika` (LTV/CAC/payback/ margins), price the created value with `vrednostno-cenovanje`, model recurring/ expansion mechanics with `revops`. If `xlsx` skill is available, render a real sheet; otherwise a markdown table with every formula shown.
- **04-moat-and-defensibility.md** — for each proprietary agent/automation in `artifacts.json`: (a) the data loop — signal captured, store, feedback path, cycles accrued; (b) the switching cost — how deep it is embedded in the workflow; (c) replicate-time — argued from loop age + integration depth, not asserted. Then the AI-specific defenses diligence will demand: model-provider redundancy, gross-margin sensitivity under a provider price shock, human-authority line (§5). Critique with `advisor-christensen` (job/switch) + `advisor-kim-mauborgne` (axis).
- **05-data-room-index.md** — the 8-category diligence checklist (corporate, financial, legal, IP, team, product/tech, cap table, tax; §6) mapped to what this build can already fill vs. what lan/the founder must supply. This is the bridge to a real raise — it tells the founder exactly what is still missing.
- **evidence/** — a copy or symlink of every artifact a number points to, so a diligence reader can trace any claim to its source in one hop.

Gate 4 — green criteria

Gate 4 writes `.protocol/gates/gate-4.json`. It is green ONLY when:

1. **investor-package/** exists and contains all six numbered files above (00–05) plus a populated `evidence/` directory (an empty file is a RED — the file must resolve its section against real artifacts; an empty `evidence/` is a RED because no number can then trace to source).
2. **The rubric has zero RED cells** — every criterion carries a real number or an artifact path, none carries an adjective, and none is `N0-SOURCE`.
3. **Every number traces to evidence/** — pick any figure in the dashboard or the economics model and it links to a source artifact in one hop. A number with no traceable source is treated as a RED cell.
4. **All seven advisor lenses have been run** over the filled rubric and their verdicts recorded; a lens that killed a criterion means that criterion is not green until fixed.
5. **The thesis is one sentence** and every bracket in it is backed by a GREEN cell.
6. **G3 is green** on 100% of registered artifacts (Gate 4 refuses to open on an unhardened build — an investor package around insecure ops is a discount magnet).

While the rubric is on the draft thresholds, gate-4.json also carries a `rubricStatus: "DRAFT"` flag and the package prints the `DRAFT-RUBRIC` banner. The gate can go GREEN structurally on the draft rubric, but the SHIP annex must show the banner until Ian's docs finalize the bars — a draft-green is honest about being draft.

AWAITING IAN'S DOCS

This pillar is a **scaffold**. The following are placeholders until Ian supplies his own investor documentation, and each is explicitly marked in the files above:

- Rubric weights & house bars.** The draft bars are now **primary-source calibrated and stage/geography-tabled** (§2–§9 of the research doc) — LTV:CAC ≥ 3.6 , payback $< 15\text{mo}$, burn $2.5\text{--}3.4\times$ seed $\rightarrow \leq 1.2\times$ A, AI GM 45–65%, EU premium $\sim 20\%$. What remains for Ian's docs is the **house layer only**: the relative *weights* across the 8 criteria, which *stage* AIS packages target by default, and any criterion he adds/removes (e.g. an AIS-house “AI-involvement score” criterion, a Slovenia/EU market criterion). The *numbers* are researched; the *emphasis* is Ian's.
- The house thesis template.** The `00-thesis.md` format is a first draft. Ian's positioning of AIS-built companies (what category, what “why now”) replaces it.
- The data-room standard AIS ships.** `05-data-room-index.md` uses the generic 8-category checklist. Ian's docs may prescribe an AIS-branded data-room template and which documents AIS produces vs. the founder supplies.
- Stage calibration.** Bars differ pre-seed vs. seed vs. Series A (§1, §3). Ian's docs set which stage AIS packages target by default and how the rubric flexes.
- The pricing/valuation link.** Whether and how the investor package feeds AIS's own pricing of the engagement (`client-pricing-sheet`) — a value-capture decision only Ian makes.

When Ian's docs arrive: replace the DRAFT bars in the rubric, remove the `DRAFT-RUBRIC` banner and `rubricStatus` flag, update `knowledge/investor-criteria-research.md` with any Ian-specific criteria, and regenerate the knowledge PDF via `make-pdf`.

Reuse map (compose before you build)

Need	Reuse
Unit economics model	<code>enotna-ekonomika</code>
Price the value created	<code>vrednostno-cenovanje</code>
Recurring/expansion mechanics	<code>revops</code>
Category + one-line positioning	<code>pozicioniranje</code>
Market/TAM bottoms-up	<code>niche-research</code> , <code>competitor-profiling</code>
Adversarial rubric review	<code>advisor-skok</code> , <code>advisor-campbell</code> , <code>advisor-hormozi</code> , <code>advisor-christensen</code> , <code>advisor-dunford</code> , <code>advisor-kim-mauborgne</code> , <code>advisor-marks</code>
Security-posture evidence	pillar ⑤ G3 output, <code>cso</code>
Render the knowledge/package to PDF	<code>make-pdf</code>

Compose these. Do not re-derive an economics engine or a positioning framework by hand — they are bundled for exactly this pillar.