Investor-Readiness Criteria — Research Findings (2026)

July 17, 2026

CONTENTS

Investor-Readiness Criteria — Research Findings (2026)

> Knowledge base for Pillar ④ (Investor-Readiness). This is the researched evidence behind the DRAFT rubric in `../SKILL.md`. It becomes a knowledge PDF (via `make-pdf`) in the delivered protocol. Findings are dated to 2026 sources; numbers are public benchmarks, NOT Ian's own deal data — see the SKILL's `AWAITING IAN'S DOCS` section.

Research date: 2026-07-17. All figures below carry a source; where a claim is a range or a median, that is the source's framing, not a smoothed guess.

---

§1 — How VCs evaluate startups in 2026 (the frame)

The foundational lens is unchanged — **Team, Product, Market, Traction** — but the bars moved:

- **AI concentration raised the bar for everyone.** ~80% of venture capital is now flowing into AI, so non-AI (and weak-AI) founders face a tighter market and higher scrutiny on team quality and unit economics. [thevccorner, spectup]
- **Diligence got slower and deeper.** What used to be a one-week diligence is now one to two months: financial statements, market data, legal structure, technical audits, plus reference calls with customers, industry experts, and prior colleagues. A clean, complete data room is now a speed advantage. [thevccorner]
- **Team carries the most weight early.** At pre-seed/seed the team often outweighs the product — investors back founders they believe can navigate the unknown. For AI companies specifically, the wanted shape is **deep AI expertise + domain/ commercial strength**: a lead engineer who has built end-to-end AI systems (or published), paired with a co-founder who owns the target domain. [thevccorner]
- **"What's your moat?" is a defensibility test, not a trivia question.** A strong answer is mechanical and compounding, e.g. *"every transaction we process trains our categorization model; we have 4M transactions; a new entrant starts at zero."* [thevccorner]

Implication for the rubric: Team (criterion 1) and Moat (5/7) are the two the package must make concrete, and the Moat answer must be a loop with a number.

Sources: The VC Corner — What Top VCs Look For 2026, spectup — What Do VCs Look For, SignalX — VC Due Diligence Framework, 4Degrees — 2026 VC Due Diligence Checklist, TechCon Global — 2026 VC Playbook, AI-First.

---

§2 — What makes an AI company defensible (the moat, correctly stated)

The consensus in 2026 is a sharp correction of the naïve "we have AI / we have data" claim:

- **Model access is NOT a moat.** AI compresses time-to-build, so a flashy feature or access to a frontier model is quickly matched. VCs are actively *rethinking* what a moat is because AI collapses build time. [Forbes/Majic, Valtorian]
- **Proprietary *static* data is NOT a moat either.** "The data advantage myth": owning a pile of data alone won't save you. [Value Add VC]
- **The real moat is a proprietary DATA LOOP.** Defensibility comes from a feedback loop where *every user action generates a better training signal → improves the model → drives more usage.* Dynamic, compounding data from real workflows is the actual moat. [Valtorian, percolator]
- **Time is the meta-moat.** What unites durable moats is that they took *years of real elapsed time to accumulate* — time that cannot be parallelized. A copycat with the same model still can't buy the accrued cycles. [Forgent/Medium]
- **Workflow depth + switching cost.** The more the product becomes part of how a team operates — with human review, exceptions, and structured states — the harder it is to rip out. Deep integration into mission-critical (especially regulated) workflows creates real switching costs. [Valtorian, Codurance, Stanford Law]
- **Domain specialization beats general-purpose.** AI built for a defined use case (grid operators, architects, lawyers) embeds in ways that feel necessary and are hard to copy. [Latitude Media]
- **The 2026 shortlist of what still defends:** *workflow depth, proprietary data loops, distribution, switching costs, and execution quality* — not model access. [Valtorian]

Implication for the rubric: Criterion 5/7 must, per agent, name the loop (signal → store → feedback → cycles), the switching cost (integration depth), and argue replicate-time from loop age. This is exactly what pillars ②③ instrument.

Sources: Valtorian — AI Moats in 2026, Forbes/Majic — VCs Rethink Startup Moats, Value Add VC — The Data Advantage Myth, percolator — Building Defensible Start-ups, Forgent — Defensibility in the Age of AI (2026), Stanford Law — Defensible Moats for Vertical AI (PDF, 2026), Latitude Media — Can startups still build a moat, Codurance — Durable Moats in the AI Era.

---

§3 — Unit-economics benchmarks (2026)

The numbers that seed the DRAFT rubric's criterion 4. **These are the placeholders pending Ian's docs.**

Metric	Healthy / benchmark (2026)	Notes
LTV:CAC	≥3:1 healthy; 3:1–5:1 typical; top quartile 5:1+	B2B SaaS median **3.2:1**. By ACV: Enterprise (>$100K) 4.5:1, Mid-market ($15K–100K) 3.2:1, SMB (<$15K) 2.5:1.
CAC payback	≤12 months healthy	Median stretched to **18 months** (up from 14 in 2023, +29% in a year); 4th quartile past 24 months.
Cost to acquire $1 ARR	—	Median SaaS now spends **~$2.00 to acquire $1 of new ARR** (up 14% since 2023).
Gross margin	70–85% healthy	Traditional SaaS clears **77–81%**. WARNING: LLM-native companies drop to **~52%** because inference cost drags margin — a live risk for AI companies.
NRR (net revenue retention)	>100%	Above 100% with evidence it is *improving* is what Series B wants.
Rule of 40	growth% + profit% > 40	Rule-of-40 companies command a **129% valuation premium** in 2026 (vs 23% in 2022).
Burn multiple	<1.5	Series B expectation.
Churn	<1.5%/mo SMB; <0.75%/mo mid-market/enterprise	Series B expectation.

Implication: the ~52% LLM-native gross-margin drag is a trap for the exact companies this protocol builds. Criterion 4 and the moat doc (§5, provider price shock) must show the company's margin math survives inference cost — or explain the architecture that keeps it high (caching, smaller models, on-prem, human-in-loop volume control).

Sources: Foundry CRO — LTV:CAC Benchmarks 2026, Beancount.io — 2026 SaaS Metrics Stack, RaiseReady — SaaS Unit Economics Bible 2026, Eagle Rock CFO — SaaS Benchmarks by Stage 2026, Digital Applied — SaaS Unit Economics 2026.

---

§4 — The AI-native operating-leverage story (revenue per employee)

The single most investor-legible metric for a company this protocol builds:

- **AI-native firms lead on revenue per employee** — the defining 2026 metric. [Forbes/Baier]
- **Magnitude: $2–4M ARR per employee** for AI-native companies, roughly **4×** the traditional-SaaS benchmark; some outliers far higher (Midjourney: ~$200M revenue / ~11 people ≈ **$18M/employee**). Other sources frame it as **2–10×** traditional SaaS. [nicktalwar, RiffOn]
- **Why it happens — structural, not per-head heroics.** AI-native orgs run with ~25% fewer people, flatter hierarchies, and higher valuation per employee. ~45% of their headcount is engineering/science (vs ~36% at non-AI startups) — *not* because they hire differently per role, but because the work that used to sit in sales, ops, finance, and admin now happens *in the product*. [Forbes/Sviokla]
- **The new thesis category:** sub-150 headcount + $100M+ ARR is "a genuinely new category of investment thesis," and the moat is explicitly *"the product loop itself"* — not patents, not static data, not exclusive distribution. [Charaka]

Implication: this is the headline of `01-metrics-dashboard.md` and `00-thesis.md`. Pillar ②'s baseline (hours/errors per process, pre-AI) → pillar ③'s automation → the revenue-per-employee and cost-out delta IS the operating-leverage proof. It is the protocol's strongest natural claim because the build literally produces it.

Sources: Forbes/Baier — AI-Native Firms Lead in Revenue Per Employee, nicktalwar — $4M Revenue Per Employee, RiffOn — 2–10× Higher Revenue Per Employee, Charaka Notes — The $5M Employee / AI-Native Revenue Density, Forbes/Sviokla — AI-Native Firms Are Flatter, Leaner, More Valuable, Valere — AI-Native Companies & Operating Leverage.

---

§5 — AI-specific technical due diligence (2026): the discount to avoid

For a company whose *whole pitch* is AI-run ops, this is where the package either earns a premium or takes a haircut. Buyers now test **four AI-vulnerability axes:** *model dependency, data moat, agentic substitution, talent concentration.* [Valutico]

- **Governance is the overlooked gap, not the model.** "The AI risk in a portfolio company is not in the model. It is in the human governance architecture that standard technical diligence never examines." Buyers ask for a documented **Human Authority Line** for every high-risk AI system, and which decisions are explicitly **non-delegable to AI.** [Falkovia]
- **Model dependency must be defended with artifacts, not assurances.** Sellers who survive this show three things: (1) a documented **model-provider redundancy plan**, (2) **gross-margin sensitivity tables under provider price shocks**, and (3) evidence the features work. [Falkovia, Valutico]
- **The haircut is real and non-negotiable.** Unaddressed regulatory, privacy, and technical AI risk draws a **15–30% valuation discount**, applied *before* other discounts stack. "Buyers do not negotiate vulnerability findings; they document them and reduce the price." [Valutico]
- **Regulatory clock — EU AI Act.** High-risk provisions begin applying in 2026, with major requirements activating around **August 2026** — directly relevant to a Slovenia/EU company (AIS's market). [Security Boulevard, governance-intelligence]
- **Operational risk sources:** data quality, model behavior, prompts, and third-party dependencies; limited visibility into training/decisions → compliance exposure, vendor lock-in, instability. [Security Boulevard, accuknox]

Implication: Pillar ⑤ (G3) is not just internal hygiene — its evidence (least-privilege, RLS proofs, secret scanning, plus a Human Authority Line and a model-provider redundancy plan) is a *valuation lever*. Criterion 8 (Security posture) and the moat doc's AI-defenses section exist to convert §5 from a –15–30% haircut into a diligence-fast-track. The package must include the authority line and the margin-sensitivity-under-price-shock table by name.

Sources: Security Boulevard — AI Due Diligence Checklist 2026, Falkovia — The AI Risk Your Due Diligence Isn't Catching, Valutico — AI Vulnerability in M&A Due Diligence: 2026 Buyer's Framework, OneTrust — Responsible AI in 2026, AccuKnox — AI Security & Governance Guide 2026, Governance Intelligence — AI & Compliance 2026.

---

§6 — What an investor-ready data room contains (2026)

Criterion structure and the `05-data-room-index.md` checklist come from here.

- **Volume:** a well-prepared room is **50–70 documents** across **8 categories**: corporate, financial, legal, IP, team, product/tech, cap table, tax. **Seed skews 40–50 docs; Series A needs 60–70.** [Peony, Papermark]
- **By category (essentials):**
  - *Corporate/Legal:* Certificate of Incorporation, bylaws, amendments, Certificate of Good Standing, board minutes & consents, stockholder/voting agreements.
  - *Financials:* seed → pitch deck + basic financials + cap table + incorporation. Series A adds GAAP statements, customer contracts, **12–24 months** of income statement / balance sheet / cash-flow, revenue-stream and cost breakdown (so investors see margins and burn), and a **12–18 month** forecast with clear assumptions. *Not* a five-year fantasy model — Series A cares about near-term use of funds. [Papermark, ascentcfo]
  - *Cap table:* founders, option pool, investors, ownership %; clean Carta export if used.
  - *Team:* org chart, key-leader bios, roles this round funds, employment/contractor agreements, and **IP assignments (PIIAs)** signed by every employee & contractor (confirming the company owns the work).
  - *Product/Tech:* product summary & roadmap, current stage (MVP/beta/launched), what this round funds.
  - *Pitch:* the core narrative — problem, solution, market, business model, competition, team — crisp and data-backed.
- **Hygiene that signals competence:** update monthly; version files with dates (`2026-03_Financial_Model_v4.xlsx`, never `Financial Model FINAL.xlsx`); include a master doc explaining the structure. [Papermark, Visible]

Implication: `05-data-room-index.md` maps these 8 categories to *what this build already produces* (product/tech summary, architecture narrative, metrics, security evidence) vs. *what the founder/Ian must supply* (corporate, legal, cap table, tax). That split is the honest boundary of what the protocol can and cannot generate.

Sources: Peony — Startup Data Room Checklist (60-doc standard), Papermark — Ultimate Startup Data Room Checklist 2026, Haven — Series A Data Room Checklist, Visible.vc — Startup Data Room, AscentCFO — Fundraising Data Room Checklist.

---

§7 — How this maps to the DRAFT rubric

Research §	Feeds rubric criterion	Key DRAFT bar it sets
§1 (frame)	1 Team, whole rubric	team = named operator + built system as proof
§2 (moat)	5 Moat, 7 AI-defensibility	≥1 agent with described compounding loop
§3 (economics)	4 Unit economics	LTV:CAC ≥3:1, GM ≥70%, payback ≤12mo
§4 (leverage)	3 Traction, 4, 6 Scalability	revenue/employee headline; volume w/ flat headcount
§5 (AI diligence)	8 Security posture, 7	G3 green on 100% artifacts + authority line + provider redundancy
§6 (data room)	05-data-room-index.md	8 categories mapped: built vs. founder-supplied

Reminder of the boundary: every number in §3, §4 is a *2026 public benchmark*, a starting placeholder. Ian's own investor documentation replaces these with his deal-tested bars, weights, stage calibration, and any AIS-specific criteria — at which point the `DRAFT-RUBRIC` banner comes off and this doc is re-issued as a knowledge PDF via `make-pdf`.